

Tutorial 05 - Midterm, UDFS

CS2040S Semester 1 2023/2024

Set real display name



<https://pollev.com/rezwanarefin430>

Midterm Review

Midterm Question 1: Sorting Problem

- Given array of n integers $A[1 \dots n]$. Find place of each $A[i]$ in sorted order.
- Constraints:
 - $n \leq 200000$.
 - $-100000 \leq A[i] \leq 99999$.
 - All the $A[i]$ are distinct.
- Example: $A = [5, -1, 4]$. Output = $[3, 1, 2]$.
 - Let's see the examples from Midterm.

Midterm Question 1: Sorting Problem

- **$O(n \log n)$ Solution:**
 - Sort **A** and store in **B**.
 - For each **A[i]**, binary search to find where it appears in **B**.

Midterm Question 1: Sorting Problem

- **$O(n \log n)$ Solution:**
 - Sort **A** and store in **B**.
 - For each **A[i]**, binary search to find where it appears in **B**.
- **$O(n)$ Solution:**
 - Use counting sort to sort **A**.
 - Offset each element by **100000**. So $0 \leq A[i] < 200000$.
 - Offsetting doesn't affect answer.
 - Note that you can deduce the output in the last step of counting sort.

Midterm Question 1: Sorting Problem

- **$O(n \log n)$ Solution:**
 - Sort **A** and store in **B**.
 - For each **A[i]**, binary search to find where it appears in **B**.
- **$O(n)$ Solution:**
 - Use counting sort to sort **A**.
 - Offset each element by **100000**. So $0 \leq A[i] < 200000$.
 - Offsetting doesn't affect answer.
 - Note that you can deduce the output in the last step of counting sort.

Midterm Question 2: Stack with `findMin()`

- Implement a Stack that supports standard `peek()`, `push()`, `pop()`, and additionally `findMin()` operation all in $O(1)$.

Midterm Question 2: Stack with `findMin()`




- Idea is to maintain a second stack that contains all the elements that may be returned by a future `findMin()` question.

Midterm Question 2: Stack with `findMin()`

- Idea is to maintain a second stack that contains all the elements that may be returned by a future `findMin()` question.
- The second stack should not contain any element that will never be answer to a `findMin()` query.


Midterm Question 2: Stack with `findMin()`

- Idea is to maintain a second stack that contains all the elements that may be returned by a future `findMin()` question.
- The second stack should not contain any element that will never be answer to a `findMin()` query.



Midterm Question 2: Stack with findMin()


- Suppose X is being pushed in the Main stack.
- If X is bigger than top element in Min stack, it will never be answer of a findMin() query.
- Otherwise it will be, so we push it in Min stack as well.






Midterm Question 2: Stack with `findMin()`

- Notice that top of the Min stack is the answer to `findMin()` now.
- What is the meaning of 2nd top element of Min stack?




Midterm Question 2: Stack with `findMin()`

- Notice that top of the Min stack is the answer to `findMin()` now.
- What is the meaning of 2nd top element of Min stack?
 - Second minimum element.



Midterm Question 2: Stack with `findMin()`

- Therefore the popping algorithm should be as follows:
 - If top of Main stack and top of Min stack is same, pop them both.
 - Else pop only Main stack.



Midterm Question 3: Adding n integers

- Given n integers. Cost of adding A and B is $(A + B)$. Find minimum cost to sum all the integers.

Midterm Question 3: Adding n integers

- Given n integers. Cost of adding A and B is $(A + B)$. Find minimum cost to sum all the integers.
- Intuition is:
 - We will always perform exactly $n - 1$ additions.

Midterm Question 3: Adding n integers

- Given n integers. Cost of adding A and B is $(A + B)$. Find minimum cost to sum all the integers.
- Intuition is:
 - We will always perform exactly $n - 1$ additions.
 - We will try to minimize cost of each of the additions.
 - How to minimize cost of the current addition?

Midterm Question 3: Adding n integers

- Given n integers. Cost of adding A and B is $(A + B)$. Find minimum cost to sum all the integers.
- Intuition is:
 - We will always perform exactly $n - 1$ additions.
 - We will try to minimize cost of each of the additions.
 - How to minimize cost of the current addition?
 - Take the smallest two A and B available and add them.

Midterm Question 3: Adding n integers

- Given n integers. Cost of adding A and B is $(A + B)$. Find minimum cost to sum all the integers.
- Intuition is:
 - We will always perform exactly $n - 1$ additions.
 - We will try to minimize cost of each of the additions.
 - How to minimize cost of the current addition?
 - Take the smallest two A and B available and add them.
 - Techniques for proving that this eventually gives the smallest possible cost will be taught in later courses (eg. CS3230).
 - You can figure this out from doing the sample test cases.
 - Let's see examples in the Midterm.

Midterm Question 3: Adding n integers

- So our algorithm is this:
 - Find and remove two smallest integers **A** and **B** from the array.
 - **Cost += A + B.**
 - Insert **A + B** back into the array.

Midterm Question 3: Adding n integers

- So our algorithm is this:
 - Find and remove two smallest integers **A** and **B** from the array.
 - **Cost += A + B.**
 - Insert **A + B** back into the array.
- To make this fast, use a priority queue to store the array.
- Complexity: **$O(n \log n)$.**

Questions?

Union Find

Union Find: The Problem

- We have n sets initially: $\{1\}, \{2\}, \{3\}, \dots, \{n\}$.
- Need to support these two type of queries:
 - Merge the sets containing u and v .
 - If u and v in the same set?

Union Find: Example

- Initial Set:
 - $\{1\}, \{2\}, \{3\}, \{4\}, \{5\}$
- Union(1, 3)
 - $\{1, 3\}, \{2\}, \{4\}, \{5\}$
- Union(2, 5)
 - $\{1, 3\}, \{2, 5\}, \{4\}$
- IsSameSet(2, 5) = False
- IsSameSet(1, 3) = True

Union Find: The Data Structure


- We will represent each set as a tree.
 - Note that the structure of the tree has nothing to do with the semantics of the operations we are supporting.
- For our algorithm, we will also elect a leader in each set. The leader will be root of the tree.
- Then merging two sets is equivalent to attaching leader of one set as a child of leader of another set.

Union Find: Example

Initial State




Union(1, 3)




Union Find: Example

Previous State




Union(2, 5)




Union Find: Example

Previous State



Union(3, 5)



Union Find: Find(u)

- Finding the leader set is a critical operation.
- Let's define Find(u) operation which returns the leader of u's set.

Union Find: Find(u)

- Finding the leader set is a critical operation.
- Let's define Find(u) operation which returns the leader of u 's set.
- Implementation:
 - Define $p[u]$ = parent of u in the tree. If $p[u] = u$, then u itself is a leader.
 - Initially everyone is a leader, so $p[i] = i$ for all i .



Union Find: Find(u)

- Finding the leader set is a critical operation.
- Let's define Find(u) operation which returns the leader of u 's set.
- Implementation:
 - Define $p[u]$ = parent of u in the tree. If $p[u] = u$, then u itself is a leader.
 - Initially everyone is a leader, so $p[i] = i$ for all i .
 - When leader u becomes subordinate of v , we will change $p[v] = u$.
 - How to implement Find(u) using this $p[]$ array?

Union Find: Find(u) Implementation

- Keep following $p[u]$ pointers until we find a leader.

```
public static int Find(int u) {  
    while (p[u] != u) u = p[u];  
    return u;  
}
```



Union Find: `IsSameSet(u, v)` Implementation

- How to check if two element are in the same set?

Union Find: IsSameSet(u, v) Implementation

- How to check if two element are in the same set?
- They must have the same leader.

```
public static boolean IsSameSet(int u, int v) {  
    return Find(u) == Find(v);  
}
```


Union Find: Union(u, v) Implementation

- Find leaders of set containing u and v .
- If they are the same, then nothing to do.
- Otherwise, attach one of them as child of another.

```
public static void Union(int u, int v) {  
    u = Find(u);  
    v = Find(v);  
    if (u == v) return;  
    p[u] = v;  
}
```

Union Find: Complexity


- Current implementation complexity is $O(n)$.
- Suppose we do $\text{Union}(1, 2)$, $\text{Union}(2, 3)$, $\text{Union}(3, 4)$, ..., $\text{Union}(n-1, n)$.
- Then we will have the following tree:




- Now if we do $\text{IsSameSet}(1, 2)$, both the calls $\text{Find}(1)$ and $\text{Find}(2)$ will take $O(n)$ to find the leader.

Union Find: Path Compression

- Note that we don't care about the structure of the tree, as long as all elements in same set are in the same tree.
- So, whenever we call $\text{Find}(u)$, we will compress the path from u to root.



- $\text{Find}(1)$ call should return n and change the structure as follows:



Union Find: Path Compression Implementation

```
public static int Find(int u) {  
    if (p[u] == u) {  
        return u;  
    } else {  
        p[u] = Find(p[u]);  
        return p[u];  
    }  
}
```


- Complexity is still worst case **$O(n)$** .
- But it can be proven that over **n** calls to Find, the total complexity cannot exceed **$O(n \log n)$** .
- So we say that the complexity is amortized **$O(\log n)$** .

Union Find: Union by Rank


- (Assume we don't have path compression)
- In $\text{Union}(u, v)$, we blindly attached u 's leader as child of v 's leader.
- If u 's tree has height 5, and v 's tree has height 4:
 - Attaching u 's tree to v 's tree results in a tree of height 6.
 - Attaching v 's tree to u 's tree results in tree of height 5.
- Attaching lower height tree to higher height tree results in a lower height.

Union Find: Union by Rank


Union(2, 3)



Attach 1 to 3



Attach 1 to 3



Union Find: Union by Rank Implementation

- Define $\text{rank}[u]$ = height of subtree rooted at u .
 - Initially everyone has rank 0.
- When merging two trees, only the leader's rank may change. Update accordingly.


```
public static void Union(int u, int v) {  
    u = Find(u);  
    v = Find(v);  
    if (u == v) return;  
    if (rank[u] > rank[v]) swap(u, v); // ensure rank[u] <= rank[v]  
    p[u] = v; // v is the new leader  
    if (rank[u] == rank[v]) ++rank[v]; // rank increases only if before merging both tree had same rank  
}
```

Union Find: Union by Rank Complexity


- With only Union by Rank, every operation is *worst case* $O(\log n)$.
- If Union by Rank and Path Compression both are used, then every operation is amortized **$O(\alpha(n))$** where α is the inverse ackermann function.
 - Ackermann function is VERY fast growing function.
 - Its inverse for any realistic value of n is < 4 .
- Note that here amortized **$O(1)$** is special. The total cost over n operation is **$O(n)$** . Normal amortized complexity would allow one single operation to be **$O(n)$** . But here Union by Rank guarantees that one single operation doesn't exceed **$O(\log n)$** .

Union Find: Union by Rank + Path Compression

- After Path Compression, the rank values may *not* correspond to height of the current structure. They will represent height of the tree assuming the tree was never compressed.



Find(1)



Tutorial Questions



Question 3: NumDisjointSets()

- How to quickly query number of disjoint sets?

Question 3: NumDisjointSets()

- How to quickly query number of disjoint sets?
- Keep a counter. Initially there are n disjoint sets.
- Whenever two sets are merged, decrease the counter by 1.
- NumDisjoinSets just returns that counter in $O(1)$.



Question 4: `SizeOfSet(u)`

- How to augment the data structure to query size of set containing u efficiently?



Question 4: `SizeOfSet(u)`

- How to augment the data structure to query size of set containing `u` efficiently?
- Similar to rank, define `size[u]` = size of set containing `u`.
 - This value will only be correct for leaders.
- Initially everyone is a leader, and `size[i] = 1` for all `i`.
- When attaching leader `u` to leader `v`, do `size[v] += size[u]`.
- `SizeOfSet(u)` should return `size[Find(u)]`.

Break Attendance Questions

[https://visualgo.net/training?diff=Medium&
n=5&tl=5&module=ufds](https://visualgo.net/training?diff=Medium&n=5&tl=5&module=ufds)

PS4 Discussion

PS4A: /swaptosort

- Given a reverse sorted array **A**.
- You have a list of **(a, b)** pairs.
- Pair **(a, b)** means that you are allowed to **swap(A[a], A[b])**.
- Can you sort the array using only allowed operations?

PS4A: /swaptosort

- Think like the midterm question.
- Where should **A[i]** go in the final array?

PS4A: /swaptosort

- Think like the midterm question.
- Where should **A[i]** go in the final array?
 - **$n - i + 1$.**
- If you need to swap indices a and b, there must be a sequence of valid swaps: **(a, x), (x, y), (y, z), ..., (*, w), (w, b)**.
- Which data structure we've learnt let us query if a and b are reachable like this?

PS4B: [/kaploeb](#)

- A simulation question.
- Use a hash table to keep track of timings, since the ids are too big.
- Parse the floating point number carefully.

Hand-On

Hands-On: /speedrun

- Given some tasks, i -th of them needs the time period $[l_i, r_i]$.
- Two tasks cannot be active at the same time.
 - Cannot do both $[1, 3]$ and $[2, 5]$.
 - Can do both $[1, 3]$ and $[3, 5]$.
- Can you do at least **G** tasks?

Hands-On: /speedrun

- Should you choose the first task you can start?
 - No! What if it takes too long?
 - Example: [1, 24000], [2, 2], [3, 3], [4, 4],
- But you can prove that choosing the task that ends first always lets you finish maximum number of tasks.
 - If choosing the task **T** that ends first was not optimal, then there is another more optimal solution
 - But in that optimal solution, you can ignore the first task and do **T** instead.
 - So choosing **T** first is also as optimal as the other solution.

Hands-On: /annoyedcoworkers

- You have c coworkers.
- Coworker i initially has a_i annoyance level.
- If you ask coworker i a question, the annoyance level increases by d_i .
- You need to ask h questions.
- Minimize the maximum annoyance level of any coworker after asking h questions.

Hands-On: /annoyedcoworkers

- Should you ask the least annoyed person first?

Hands-On: /annoyedcoworkers

- Should you ask the least annoyed person first?
 - No! What if his temper is infinite?
 - Remember you are trying to minimize the resulting max annoyance level.

Hands-On: /annoyedcoworkers

- Should you ask the least annoyed person first?
 - No! What if his temper is infinite?
 - Remember you are trying to minimize the resulting max annoyance level.
- Choose the person whose resulting annoyance level will be the smallest.

Hands-On: /annoyedcoworkers

- Should you ask the least annoyed person first?
 - No! What if his temper is infinite?
 - Remember you are trying to minimize the resulting max annoyance level.
- Choose the person whose resulting annoyance level will be the smallest.
- Use a priority queue to store the persons, sorted by their resulting annoyance level if you were to ask them a question.

Thank You!

Anonymous Feedback:

<https://forms.gle/MkETeXdUT53Vhh896>